

## ISMLA 17/18: uimaFIT Exercises

1. Port your OpenNLP tokenizer AE from UIMA to uimaFIT (create a new class that partly reuses the code from the existing component). Adapt the parent class, the configuration parameters, and the methods for iterating over annotations.
2. Wrap the OpenNLP POS Tagger (<https://opennlp.apache.org/docs/1.8.3/manual/opennlp.html#tools.postagger>) as a uimaFIT annotator and test it by creating an ad-hoc JCas and calling the AE's `process` function on the JCas. A POS tagger expects sentence-segmented and tokenized input.
3. Port your Collection Reader and CAS Consumer as uimaFIT components and run them via `SimplePipeline.runPipeline`. If time is too short, you can also use the ported versions of the reader and consumer provided on Moodle, understand the code, and assemble them in a pipeline together with your ported AEs.  
The code should be written in a static function that takes at least two arguments: *input directory* and *output directory*.